

# The K Project

## Introduction

LSE Team

EPITA

septembre 27, 2017



Figure: K running 'skate'

The K Project

LSE Team

Introduction

Debugging

x86  
Architecture

Serial

Conclusion

- Serial driver
- Segmentation
- Events
- Keyboard
- Timer
- Syscalls
- VGA driver
- File System
- Binary loading
- Bonus: Sound driver, Console driver, ...

The K Project

LSE Team

Introduction

Debugging

x86  
Architecture

Serial

Conclusion

- Basic serial driver
- Segmentation initialization

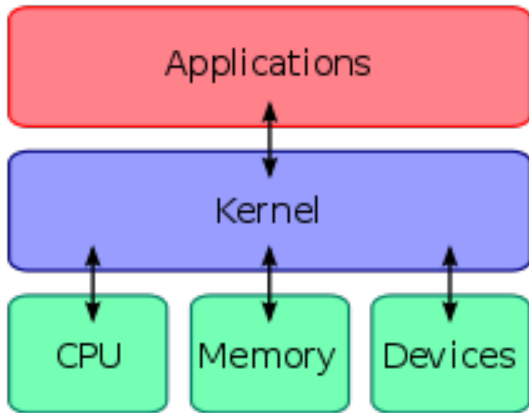


Figure: Operating system layout

The K Project

LSE Team

Introduction

Debugging

x86  
Architecture

Serial

Conclusion

```
git clone git://git.lse.epita.fr/k.git
```

The K Project

LSE Team

Introduction

Debugging

x86

Architecture

Serial

Conclusion

```
|-- config.mk
|-- k
|   |-- compiler.h
|   |-- crt0.S
|   |-- elf.h
|   |-- include
|       |-- k
|   |-- io.h
|   |-- k.lds
|   |-- k.c
|   |-- libvga.c
|   |-- libvga.h
|   |-- Makefile
|   |-- multiboot.h
|-- Makefile
|-- roms
|-- libs
|   |-- libc
|   |-- libk
|-- tools
|   |-- create-iso.sh
|   |-- mkkfs
|   |-- mkksf
```

```
#include "multiboot.h"
#include "kstd.h"

void      k_main(unsigned long      magic,
                    multiboot_info_t* info)
{
    (void) magic;
    (void) info;

    unsigned int i = 0;
    char star[4] = { '|', '/', '-', '\\' };
    char* fb = (void*) 0xb8000;

    while (1)
        *fb = star[i++ % 4];
}
```



The K Project

LSE Team

Introduction

Debugging

x86

Architecture

Serial

Conclusion

- <http://k.lse.epita.fr/>
- <http://intel.com/products/processor/manuals/>

## Launch your K

```
$ qemu-system-x86_64 -cdrom k.iso [ -enable-kvm ]
```

## Have QEMU wait for your debugger to hook:

- Add “-s -S” to QEMU options

## Launch gdb and hook to QEMU

```
$> gdb k/k  
$(gdb)> target remote localhost:1234
```

## Set your breakpoint and continue

Don't forget to build with debug options !

The K Project

LSE Team

Introduction

Debugging

x86  
Architecture

Serial

Conclusion

- General purpose registers
- Segment registers
- Flags
- Control & Memory registers
- Tons of others (XMM0-7...)

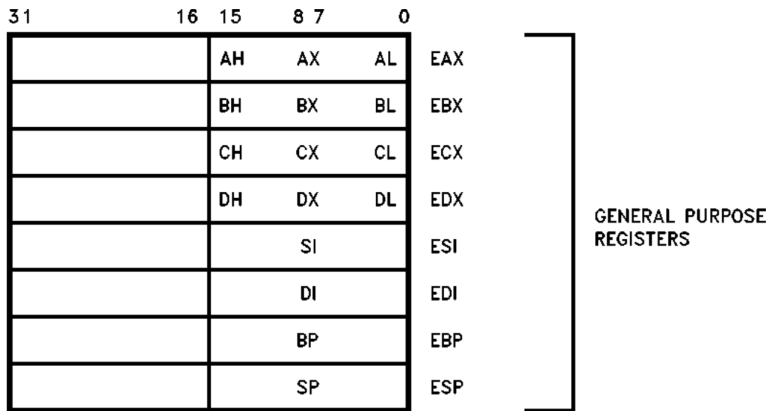


Figure: General purpose registers layout

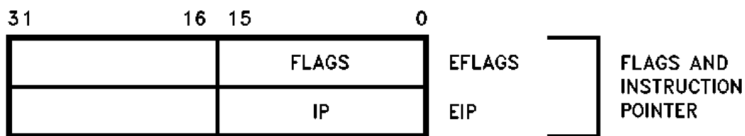


Figure: EIP/IP and EFLAGS/FLAGS

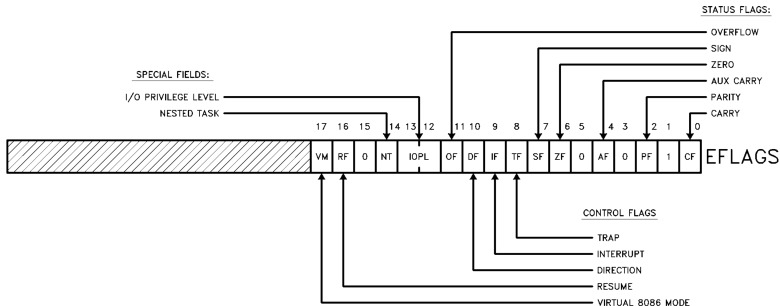


Figure: Flags layout

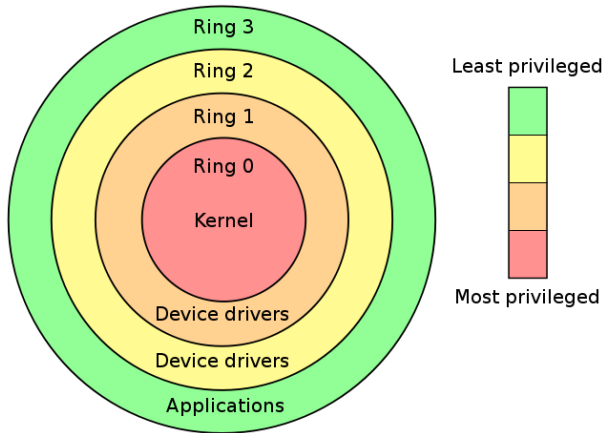


Figure: x86 privileges rings

## C declaration:

```
pushl %eax ; arg3
pushl %ebx ; arg2
pushl %ecx ; arg1
call  foo ; foo(arg1, arg2, arg3)
```

## Think of call as:

```
pushl %eip
%eip = ADDRESS
```

## Think of ret as:

```
popl %eip
```



## Asm exemple for sum(int, int)

sum:

```
pushl %ebp
movl %esp, %ebp
movl 8(%ebp), %eax ; put first arg in %eax
addl 12(%ebp), %eax ; add second arg to %eax
movl %ebp, %esp
popl %ebp
ret
```

## Basic syntax:

```
__asm__ ("movb %ah, (%ebx)");
```

## Tell GCC/GAS not to optimize your code:

```
asm volatile ("movl $0, %eax");
```

## Note

You can either write GNU keywords and specifiers with or without double underscore around them to avoid name conflicts (asm or \_\_asm\_\_, volatile or \_\_volatile\_\_).

## ASM inline template

```
__asm__("[your assembly code]"  
: output operands /* optional */  
: input operands /* optional */  
: list of clobbered /* optional */  
);
```

## Different output/input constraints:

<http://gcc.gnu.org/onlinedocs/gcc/Constraints.html>

"m" : memory operand

"r" : register operand

## Constraint modifiers

<https://gcc.gnu.org/onlinedocs/gcc/Modifiers.html#Modifiers>

- "=" : Write Only

- "+" : Read/Write

## Different clobbers

memory

[register names]

```
asm volatile("outb %0, %1\n\t"  
             : /* No output */  
             : "a" (val), "d" (port));
```

```
asm volatile("inb %1, %0\n\t"  
             : "&a" (res)  
             : "d" (port));
```

Intel Code	AT&T Code
<code>mov eax,1</code>	<code>mov \$1,%eax</code>
<code>int 80h</code>	<code>int \$0x80</code>
<code>mov ebx,eax</code>	<code>mov %eax,%ebx</code>
<code>mov eax,[ebx+3]</code>	<code>mov 3(%ebx),%eax</code>
<code>mov eax,[ebx+20h]</code>	<code>mov 0x20(%ebx),%eax</code>
<code>lea eax,[ebx+ecx]</code>	<code>lea (%ebx, %ecx),%eax</code>

```
struct {  
    unsigned char field_a : 1; // max value is 0b1  
    unsigned char field_b : 2; // max value is 0b11  
    unsigned char field_c : 5; // max value is 0x1F  
} bitfields;
```

## Note

`sizeof(bitfields)` is equal to `sizeof(unsigned char)`.

## Not Packed

```
struct {  
    unsigned char    a; // aligns with 3 bytes  
    unsigned int     b; // aligned  
    unsigned char    c; // aligns with 3 bytes  
} foo;
```

## Note

`sizeof(foo)` gives 12 (1 + 3 padding + 4 + 1 + 3 padding).



## Packed

```
struct {  
    unsigned char    a;  
    unsigned int     b;  
    unsigned char    c;  
}__attribute__((packed)) bar;
```

## Note

sizeof(bar) gives 6, struct is memory packed and no padding is inserted.

`write()`

```
int write(const char *buf, size_t count);
```

Note

`write()` sends to COM1

Note

`printf()` is available in your kernel and uses `write()`

The K Project

LSE Team

Introduction

Debugging

x86  
Architecture

Serial

Conclusion

- Separated adress space
- $2^{16}$  adresses
- in/out x86 instructions family
- Serial/PIC/PIT/Keyboard

- COM1 + 3: (8 bits length) | (No parity)
- COM1 + 2: (FIFO) | (Interrupt trigger level 14 bytes) | (Clear transmit FIFO) | (Clear receive FIFO)
- COM1 + 1: Enable Transmitter Holding Register Empty Interrupt

- A line ends with `\r\n`
- You can redirect the serial output with:  
`qemu-system-x86_64 [...] -serial stdio`

## Serial Registers:

Eight I/O bytes are used for each UART to access its registers.

The following table shows, where each register can be found.

The base address used in the table is the lowest I/O port number assigned.

The switch bit DLAB can be found in the line control register LCR as bit 7 at I/O address  $\text{base} + 3$ .

UART register to port conversion table

I/O port	DLAB = 0		DLAB = 1	
	Read	Write	Read	Write
base	<b>RBR</b> receiver buffer	<b>THR</b> transmitter holding	<b>DLL</b> divisor latch LSB	
base + 1	<b>IER</b> interrupt enable	<b>IER</b> interrupt enable	<b>DLM</b> divisor latch MSB	
base + 2	<b>IIR</b> interrupt identification	<b>FCR</b> FIFO control	<b>IIR</b> interrupt identification	<b>FCR</b> FIFO control
base + 3	<b>LCR</b> line control			
base + 4	<b>MCR</b> modem control			
base + 5	<b>LSR</b> line status	– factory test	<b>LSR</b> line status	– factory test
base + 6	<b>MSR</b> modem status	– not used	<b>MSR</b> modem status	– not used
base + 7	<b>SCR</b> scratch			

Figure: UART

**DLL and DLM : Divisor latch registers**

Speed (bps)	Divisor	DLL	DLM
<b>50</b>	2,304	0x00	0x09
<b>300</b>	384	0x80	0x01
<b>1,200</b>	96	0x60	0x00
<b>2,400</b>	48	0x30	0x00
<b>4,800</b>	24	0x18	0x00
<b>9,600</b>	12	0x0C	0x00
<b>19,200</b>	6	0x06	0x00
<b>38,400</b>	3	0x03	0x00
<b>57,600</b>	2	0x02	0x00
<b>115,200</b>	1	0x01	0x00

Figure: Divisor



FCR : FIFO control register

Bit	Value	Description
0	0	Disable FIFO's
	1	Enable FIFO's
1	0	—
	1	Clear receive FIFO
2	0	—
	1	Clear transmit FIFO
3	0	Select DMA mode 0
	1	Select DMA mode 1
4	0	Reserved
5	0	Reserved (8250, 16450, 16550)
	1	Enable 64 byte FIFO (16750)
Bit 7 Bit 6		Receive FIFO interrupt trigger level
6,7	0 0	1 byte
	0 1	4 bytes
	1 0	8 bytes
	1 1	14 bytes

Figure: FCR

**IER : Interrupt enable register**

Bit	Description
0	Received data available
1	Transmitter holding register empty
2	Receiver line status register change
3	Modem status register change
4	Sleep mode (16750 only)
5	Low power mode (16750 only)
6	reserved
7	reserved

Figure: IER

LCR : line control register

Bit	Value		Description	
0,1	Bit 1	Bit 0	Data word length	
	0	0	5 bits	
	0	1	6 bits	
	1	0	7 bits	
	1	1	8 bits	
2	0		1 stop bit	
	1		1.5 stop bits (5 bits word)	
			2 stop bits (6, 7 or 8 bits word)	
3,4,5	Bit 5	Bit 4	Bit 3	
	x	x	0	No parity
	0	0	1	Odd parity
	0	1	1	Even parity
	1	0	1	High parity (stick)
	1	1	1	Low parity (stick)
6	0		Break signal disabled	
	1		Break signal enabled	
7	0		DLAB : RBR, THR and IER accessible	
	1		DLAB : DLL and DLM accessible	

Figure: LCR

The K Project

LSE Team

Introduction

Debugging

x86  
Architecture

Serial

Conclusion

■ `lionel[at]lse.epita.fr` with `[K]` tag